

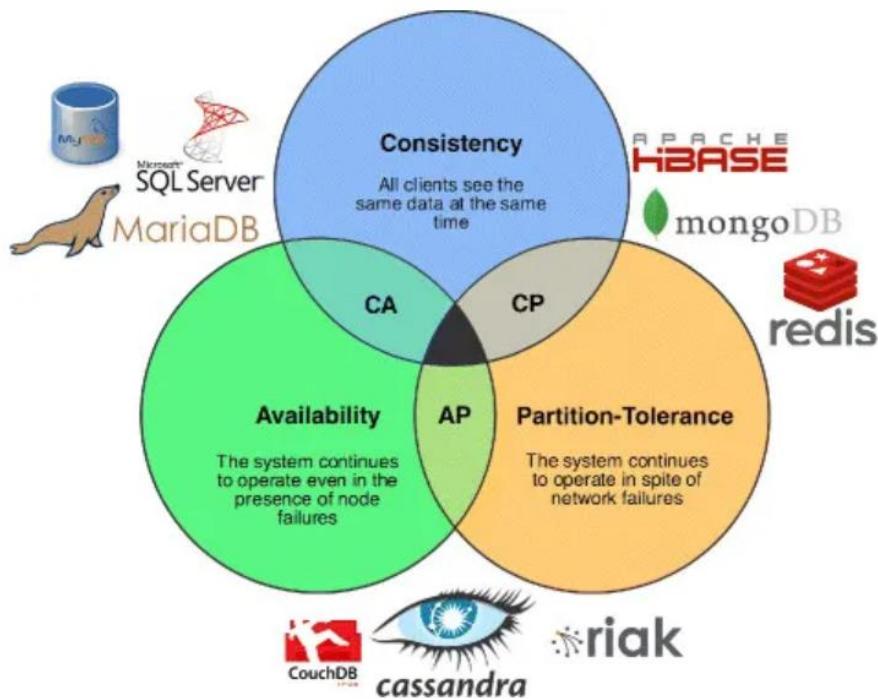
CAP Theorem

The **CAP theorem** (also known as Brewer's Theorem) states that **in a distributed database system, it is impossible to guarantee all three of the following properties simultaneously:**

Property	Meaning
C – Consistency	Every read receives the most recent write or an error. All nodes return the same data at any given time.
A – Availability	Every request (read or write) gets a response , even if some nodes have failed (but data may not be the latest).
P – Partition Tolerance	The system continues to operate even if there is a network partition (communication break) between nodes.

🔑 Rule

- **You can only guarantee TWO of the three** at the same time.
- In a real-world distributed system, **Partition Tolerance (P)** is unavoidable (network failures happen),
so the real trade-off is usually between **Consistency (C)** and **Availability (A)**.



Source : medium.com

CAP Theorem states that in a **distributed database system**, it is **impossible to guarantee all three** of the following properties at the same time:

- **Consistency (C)**
- **Availability (A)**
- **Partition Tolerance (P)**

- **Consistency**

- All nodes must show the **same data at any given time**.
- Clients always receive the **most recent write** when reading.
- Requires **instant replication** of new writes to all nodes.

- **Availability**

- The system must **always respond** to requests (read/write).
- Works even if some nodes fail.
- Requires redundancy and fault-tolerant mechanisms.

- **Partition Tolerance**

- The system continues to operate **even if network partitions occur** (nodes are split and cannot communicate).
- Uses techniques like replication, redundancy, consensus algorithms, and load balancing.

- **CA (Consistency + Availability)**

- Cannot handle network partitions well.
- Example: Traditional relational databases (MySQL in single-node setups).

- **CP (Consistency + Partition tolerance)**

- Sacrifices availability during network issues to ensure data accuracy.
- Example: MongoDB (replica sets), HBase.

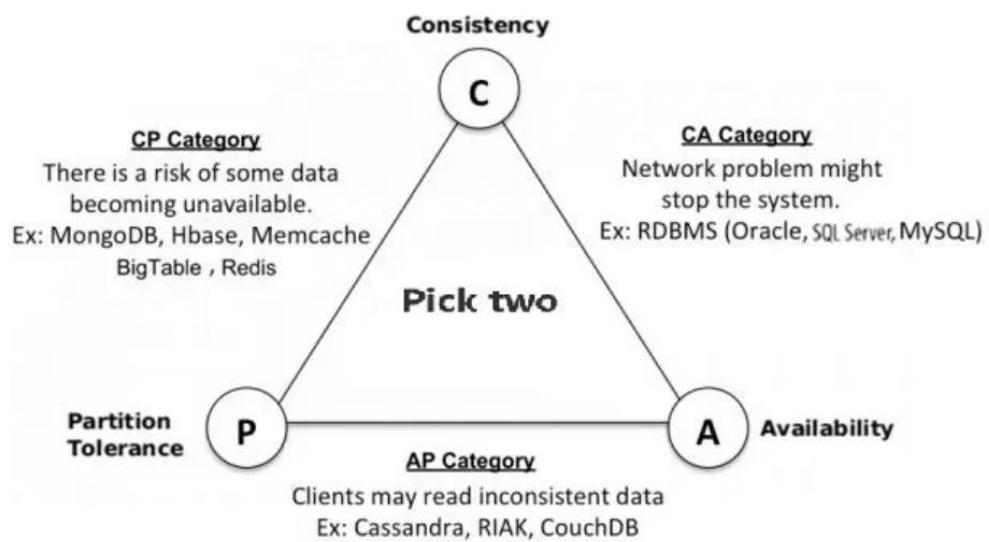
- **AP (Availability + Partition tolerance)**

- Sacrifices strong consistency to remain operational during partitions.
- Example: Cassandra, CouchDB.

Real-World Examples

Database Type	Trade-Off	Use Case
Relational (MySQL)	CA	Strong consistency required, single-node or small networks.
NoSQL (Cassandra)	AP	High availability and scalability, tolerates eventual consistency.
NewSQL (Google Spanner)	Balance CA/CP	Uses advanced methods (e.g., synchronized clocks) for global consistency.

- **Partition tolerance is non-negotiable** in distributed systems.
- Designers must **choose between stronger consistency or higher availability** based on application requirements:
 - Banking → **CP** (consistency critical).
 - Social media → **AP** (availability critical).



Source :medium.com

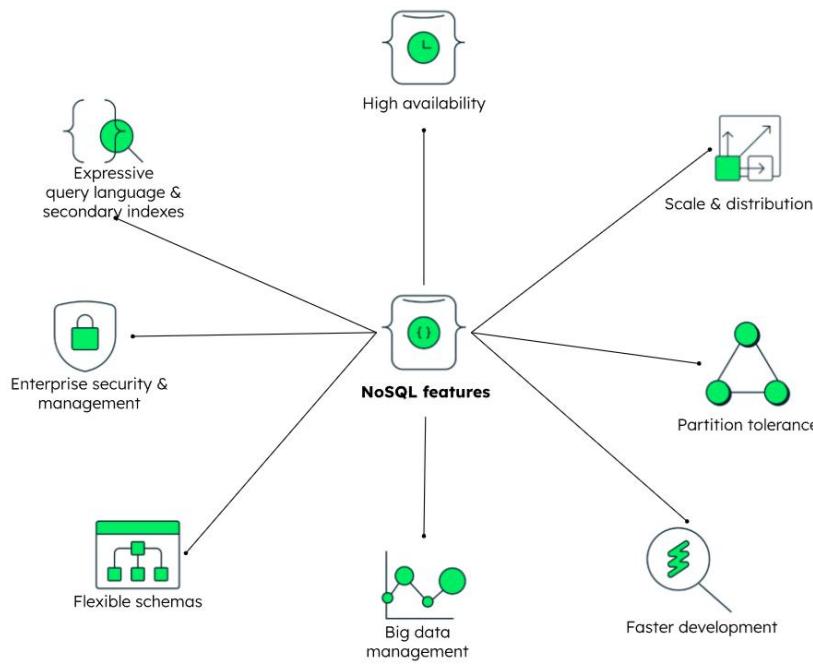
The BASE model (Basically Available, Soft State, Eventual Consistency)

- **BASE** stands for:
 - **Basically Available**
 - **Soft State**
 - **Eventual Consistency**
- It is a **consistency model** used in **NoSQL databases** (including MongoDB) that **prioritizes availability and partition tolerance** over immediate consistency.

Principles of BASE

Principle	Explanation
Basically Available	The system remains available to process requests even during network partitions or partial failures. Achieved using data distribution and replication so that some replicas are always accessible.
Soft State	The system's state may change over time even without new input , as data propagates and synchronizes across replicas. Different nodes may temporarily hold slightly different versions of the data.
Eventual Consistency	The system guarantees that all replicas will eventually converge to the same consistent state , once updates propagate and partitions are resolved. Immediate consistency is not required.

Features of NoSQL:



MongoDB follows BASE principles through several mechanisms:

- **Replication**
 - Uses **replica sets** to ensure high availability and data redundancy.
 - If one node fails, others can continue to serve read/write operations.
- **Flexible Schema**
 - Supports a **dynamic schema**, allowing documents in a collection to have different structures.
 - Supports the *soft state* concept since schema updates do not require global changes.
- **Read Preferences & Write Concerns**
 - Developers can control how reads and writes behave:
 - Read from **primary** or **secondary** nodes.
 - Configure how many replicas must acknowledge a write.
 - Balances **availability** vs. **consistency** depending on application needs.

5 Advanced Consistency Options

- Though BASE is the default approach, MongoDB can be configured to provide **stronger consistency** when needed:
 - **Multi-document ACID transactions** (introduced in later versions) provide traditional ACID guarantees for critical operations.

6 Key Takeaway

- **BASE = High Availability + Scalability + Eventual Consistency**
- MongoDB excels in environments where:
 - **24/7 availability** is critical.
 - **Large-scale distributed data** needs to be handled.
 - Temporary **inconsistencies are acceptable** (e.g., social media feeds, e-commerce catalogs).

Difference between Replication and Redundancy

Replication

Feature	Details
Definition	The process of creating and maintaining identical copies of data across multiple servers (nodes) in real time or near real time.
Goal	Ensure high availability, fault tolerance, and read scalability .
How it works	One node acts as a Primary (handles writes), and other nodes are Secondaries (continuously replicate the primary's data).
Behavior	If the primary fails, a secondary can be promoted to primary automatically (failover).
Example in MongoDB	Replica Sets – where multiple nodes store the same data and replicate updates from the primary.
Key Point	Replication actively keeps multiple live copies of data synchronized for fast recovery and read operations.

Redundancy

Feature	Details
Definition	The practice of storing extra copies of data or infrastructure (hardware, network paths, power, etc.) to provide backup in case of failure.
Goal	Ensure data durability and system reliability even during hardware or catastrophic failures.
How it works	Can involve replication but also includes backups, extra servers, spare network links, or duplicate components .
Behavior	Redundant components may remain idle until needed (e.g., backup servers or disks).

Feature	Details
Example in MongoDB	Keeping backups on different servers/clouds , RAID storage, or using sharded clusters with redundant hardware .
Key Point	Redundancy is a broader concept —it covers any duplication (not just active data replication) to prevent single points of failure.

Difference

Aspect	Replication	Redundancy
Scope	A data synchronization technique .	A design principle to add backups or extra components.
Active/Passive	Usually active , with live updates between nodes.	Can be active (hot spare) or passive (cold backup).
Purpose	Keeps multiple live, updated copies for real-time availability.	Ensures recovery options if one or more components fail.
Example	MongoDB replica sets copying data between nodes.	RAID disks, offsite backups, or secondary data centers.

Introduction to MongoDB

- MongoDB is a **NoSQL, document-oriented** database that stores data in a **flexible, JSON-like** format called **BSON** (Binary JSON).
- Unlike traditional **Relational Databases (RDBMS)**, which store data in **tables** and **rows**, MongoDB stores **documents** inside **collections**.
- It is designed for **high performance, scalability**, and **ease of development**, making it ideal for applications with **large amounts of unstructured or semi-structured data**.

💡 Key Characteristics:

- **Schema-less:** No fixed schema is required; each document can have different fields.
- **Cross-platform:** Works on Windows, Linux, macOS, and in the cloud.
- **Scalable:** Supports horizontal scaling through **sharding** (splitting data across servers).
- **Open-source:** Free to use with enterprise and cloud versions available.

🌟 Features of MongoDB

Feature	Explanation	Example
Document-Oriented	Data is stored as documents (BSON) instead of rows and columns.	{ name: "Alice", age: 25, skills: ["Python", "MongoDB"] }
Flexible Schema	Collections do not require a predefined structure; fields can differ across documents.	One document may have an <code>address</code> field, another may not.
High Availability	Replica Sets provide automatic failover and redundancy.	If the primary node goes down, a secondary takes over.
Horizontal Scalability	Supports Sharding to distribute data across multiple machines.	Splitting a large dataset across servers for faster access.
Indexing	Supports indexes (single field, compound, geospatial, text) to speed up queries.	<code>db.users.createIndex({name:1})</code>
Aggregation Framework	Performs complex data analysis operations like filtering, grouping, or transforming data.	<code>\$group, \$match, \$sort</code>
Ad-hoc Queries	Supports rich query language with dynamic filtering and sorting.	<code>db.users.find({age:{\$gt:20}})</code>
File Storage	Stores large files (images, videos) using GridFS .	
Strong Consistency	Default write operations go to the primary, ensuring up-to-date reads.	
Cross-platform & Cloud	Works with on-premise setups or cloud solutions like MongoDB Atlas .	



MongoDB Command Interface

MongoDB provides a **command-line shell** (`mongosh`) where developers can run commands to interact with the database.



MongoDB Compass

- **MongoDB Compass** is a **Graphical User Interface (GUI)** tool provided by MongoDB to visually interact with the database.

- Ideal for users who prefer **point-and-click operations** instead of using the command line.

Key Features:

Feature	Details
Visual Data Exploration	Browse databases, collections, and documents with an intuitive interface.
Query Builder	Run queries using a visual query builder (no need to type commands).
Schema Analysis	Automatically analyzes the schema of collections to detect field types and structures.
Aggregation Pipeline	Build aggregation pipelines visually to process and analyze data.
Performance Monitoring	Provides insights into query performance and database statistics.

⚡ Quick Comparison

Feature	MongoDB Command Interface (mongosh)	MongoDB Compass
Type	Command-Line (Text-based)	GUI (Visual)
Best for	Developers comfortable with scripting	Beginners, analysts, or quick visual tasks
Speed	Faster for bulk or automated tasks	Better for exploring and learning
Control	Full control of advanced operations	Limited to GUI-supported features

JSON (JavaScript Object Notation)

✓ What is JSON?

- JSON stands for **JavaScript Object Notation**.
 - It is a **text-based** data interchange format used to store and exchange data between applications.
 - Data is represented as **key-value pairs** and is **human-readable**.
 - Widely used in APIs, web applications, and configuration files.
-

💡 Characteristics of JSON

Feature	Description	Example
Text Format	Data is stored as text (UTF-8 encoding).	"name": "Alice"
Human-readable	Easy to read and write for humans.	{ "name": "Alice", "age": 25 }
Language-independent	Can be used with most programming languages (Java, Python, etc.).	{ "id": 101, "status": true }
Data Types	Supports basic types: string, number, boolean, array, object, null.	["apple", "banana", "mango"]
Lightweight	Simple structure for easy data exchange.	{ "city": "London" }

Example JSON Document

```
{
    "name": "Alice",
    "age": 25,
    "skills": ["Python", "MongoDB"],
    "isActive": true,
    "address": {
        "city": "New York",
        "zip": "10001"
    }
}
```

BSON (Binary JSON)

✓ What is BSON?

- **BSON** stands for **Binary JSON**.
- It is the **binary-encoded** version of JSON.
- MongoDB internally stores data in BSON format instead of plain JSON.
- BSON extends JSON by supporting additional data types (e.g., **Date**, **ObjectId**, **Decimal128**, **Binary Data**).

- Decimal128
 - **128-bit precision**
 - **34 decimal digits of accuracy**

Examples:

currency values like ₹123456789.123456789

Scientific calculations

💡 Characteristics of BSON

Feature	Description	Example
Binary Format	Data is stored in a binary-encoded format, making it more efficient for machines to parse.	Stored in bytes, not text.
Supports More Data Types	Includes types not present in JSON, such as Date , ObjectId , Binary , Int32 , Int64 , etc.	<code>ObjectId("507f1f77bcf86cd799439011")</code>
Efficient	Faster to encode/decode and smaller in size for certain structures (e.g., integers, binary data).	Useful for databases.
Machine-friendly	Optimized for speed and performance when storing and querying data.	Internal MongoDB storage.



Example BSON Document

A JSON document:

```
{  
    "name": "Alice",  
    "age": 25,  
    "joinDate": "2025-09-20T10:00:00Z"  
}
```

The same document in BSON might internally look like:

```
\x16\x00\x00\x00      // total document size  
\x02 name \x00 \x06\x00\x00\x00 Alice\x00  
\x10 age \x00 \x19\x00\x00\x00 // 25 as a 32-bit integer  
\x09 joinDate \x00 \x80\xca...  
\x00      // document terminator
```

(Shown as binary representation, not human-readable)

- \x02 = String type
- \x10 = Integer type
- \x09 = Date type

MongoDB understands and stores this efficiently in binary.

JSON vs. BSON

Feature	JSON	BSON
Format	Text-based (UTF-8)	Binary-encoded (machine-readable)
Readability	Human-readable	Machine-readable
Data Types	Limited (String, Number, Boolean, Array, Object, Null)	Extended (Date, ObjectId, Decimal, Binary, Int32/Int64)
Efficiency	Larger and slower to parse	Compact and faster for database operations
Usage	Data exchange (APIs, config files)	Internal MongoDB storage and queries
Example	{ "age": 25 }	Binary representation of { "age": 25 }

Why MongoDB Uses BSON

MongoDB uses BSON **internally** because:

1. **Speed:** Binary format is faster to parse and query.
 2. **Extra Data Types:** Supports ObjectId, Date, Binary, Decimal, etc.
 3. **Size Efficiency:** Smaller storage size for certain data structures.
-

⚡ Key Takeaways

- **JSON:** Ideal for communication between client and server (human-readable).
- **BSON:** Used by MongoDB for **storage**, **indexing**, and **querying** because it is faster and supports more data types.
- MongoDB accepts data as JSON when you insert (`db.collection.insertOne()`), but internally **converts it to BSON** for storage.